

Mihai Todor
28.05.2020

Integrating the Envoy gRPC API into a Dynamic Service Discovery Platform

About me

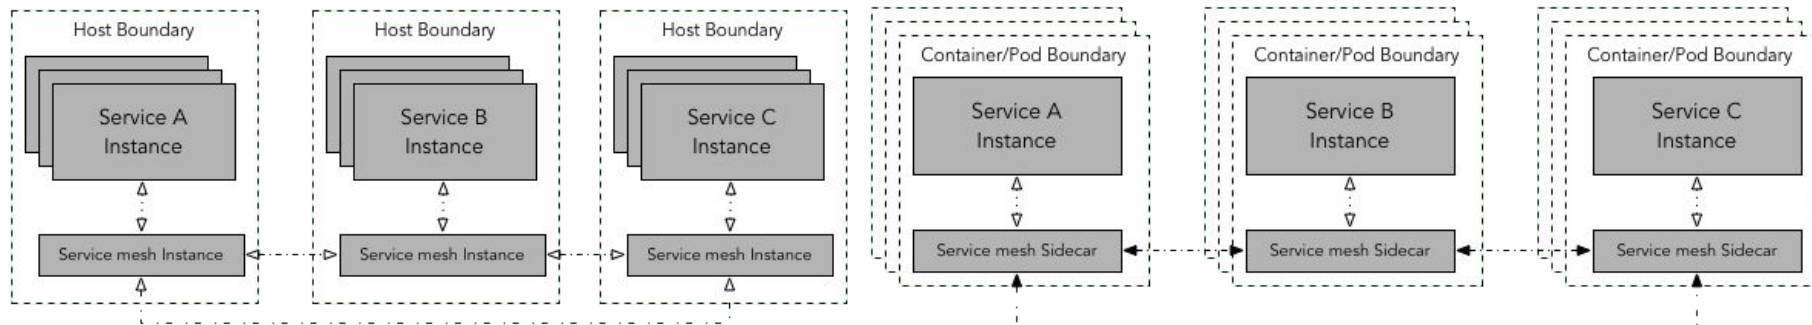
- Senior Software Engineer at [Cogito](#)
- Focused on highly-scalable distributed systems and Go
- Find me on:
 - Linkedin: <https://www.linkedin.com/in/mtodor>
 - Twitter: [@MihaiTodor](#)

Service meshes

- Infrastructure layer of an enterprise service cluster
- Handle service-to-service communication
 - Reliability
 - Security
 - Observability
 - Management
- Examples:
 - Istio <https://istio.io/>
 - Consul <https://www.consul.io/>
 - Linkerd <https://linkerd.io/>
 - Kuma <https://kuma.io/>
 - Maesh <https://containo.us/maesh/>




Service mesh deployment models



Per-host proxy deployment

Sidecar proxy deployment

Sidecar <https://github.com/Nitro/sidecar>




- Dynamic service discovery platform
- Per-host proxy via either [HAProxy](#) (through [haproxy-api](#)) or [Envoy](#)
- Docker native
- Gossip-based communication between hosts via [Memberlist](#)
- Health checks (HTTP or external)
- Has been used in production on [Apache Mesos](#) clusters

Sidecar history

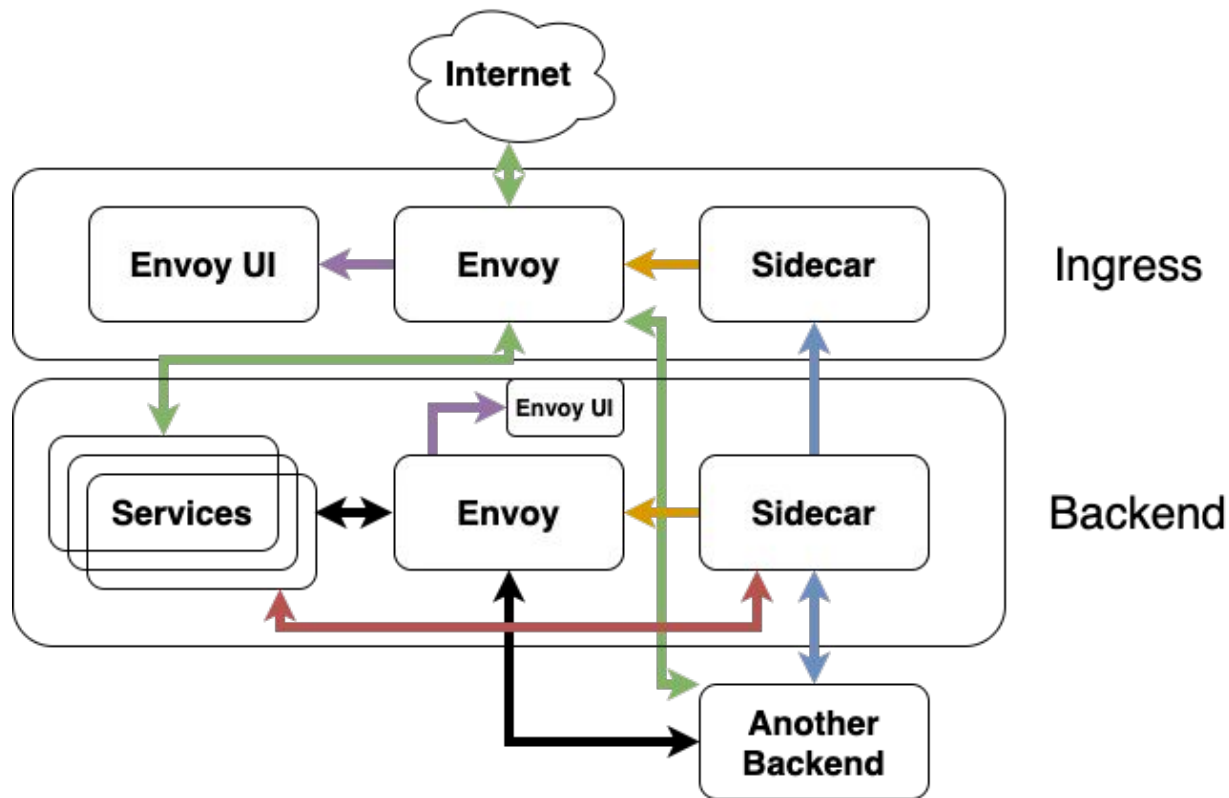
- [Karl Matthias](#) started developing it at New Relic back in [2015](#)
 - I'll have to ask him what “bosun” stands for :)
- He deployed a fully-working version in production at [Nitro](#) in 2016 on top of a Mesos cluster, initially using HAProxy as a sidecar proxy
 - Envoy support was added afterwards via the now deprecated Envoy REST API
- He continues to use it in production at [Community](#) and maintain it
- I recently [integrated](#) the Envoy [go-control-plane](#) to enable support for the Envoy gRPC API
 - I used the [Aggregated Discovery Service](#), which is part of the [xDS gRPC-based V2 API](#)

Architecture



Detail

- Ingress traffic
- Inter-service traffic
- Sidecar gossip
- Sidecar Docker discovery
- Sidecar health checks
- Sidecar -> Envoy updates
- Envoy UI updates



Demo

- Simulated cluster on local laptop using Docker-in-Docker aka [DinD](#)
- One ingress container running
 - Sidecar in listener mode
 - Envoy with static listeners and [ADS](#)
 - [Envoy UI](#)
- Three backend containers (or more) running
 - Sidecar
 - Envoy with [ADS](#)
 - [Envoy UI](#)
 - Three [WhoAmI](#) containers
 - two HTTP services
 - one TCP service

Envoy APIs

- v1 REST-JSON xDS API (deprecated and no longer supported)
- v2 xDS API (deprecated, end-of-life EOY 2020)
- Beyond...

Push vs Pull based APIs

- The Sidecar and Envoy states are designed to be eventually consistent
- For the V1 API, Envoy was configured to pull the whole state from Sidecar every 4 seconds
- For the V2 API, Sidecar checks its internal state for updates every second and pushes the updated state to Envoy if needed
 - Alternatively, it could send an update on each state update event, but this has potential issues:
 - Spurious updates
 - Potentially missed updates

State updates

- Can be expensive to send the whole state on each update in a large cluster
- We can leverage the Envoy [Incremental xDS](#) API to send partial updates
 - [Eventual consistency considerations](#)
 - Still [work in progress](#) in the Envoy go-control-plane

Envoy go-control-plane ingredients

- xDS server: github.com/envoyproxy/go-control-plane/pkg/server
- gRPC server: google.golang.org/grpc
- github.com/envoyproxy/go-control-plane/envoy/service/discovery/v2
 - [RegisterAggregatedDiscoveryServiceServer](#) to connect the gRPC and xDS servers
- Resource cache: github.com/envoyproxy/go-control-plane/pkg/cache
 - [SetSnapshot](#) instructs Envoy that it needs to fetch the updated [resources](#)
 - Sets Envoy [listeners](#), [clusters](#) and other resources
 - Requires a [new version](#)
 - The **node** parameter [needs to match](#) the value passed via `--service-node` to Envoy
 - Envoy validates clusters by default for the [HTTPConnectionManager](#)-based listeners
 - If enabled, clusters need to be added before adding listeners that depend on them;
See the [eventual consistency considerations](#)
 - Can be disabled via the `validate_clusters` parameter of the [RouteConfiguration](#)

Testing strategies

- Mock the Envoy go-control-plane... Uh-oh
- Mock Envoy itself by creating a dummy gRPC client
 - Nonces need to be [passed around correctly](#)
- Go makes it trivial to inherit all the data members and methods of a struct and override desired methods as needed
 - Enable [various assertions](#) during concurrent workflow using a [channel-based blocking state machine](#)

Thank you!

Please let me know if you have any questions 😊